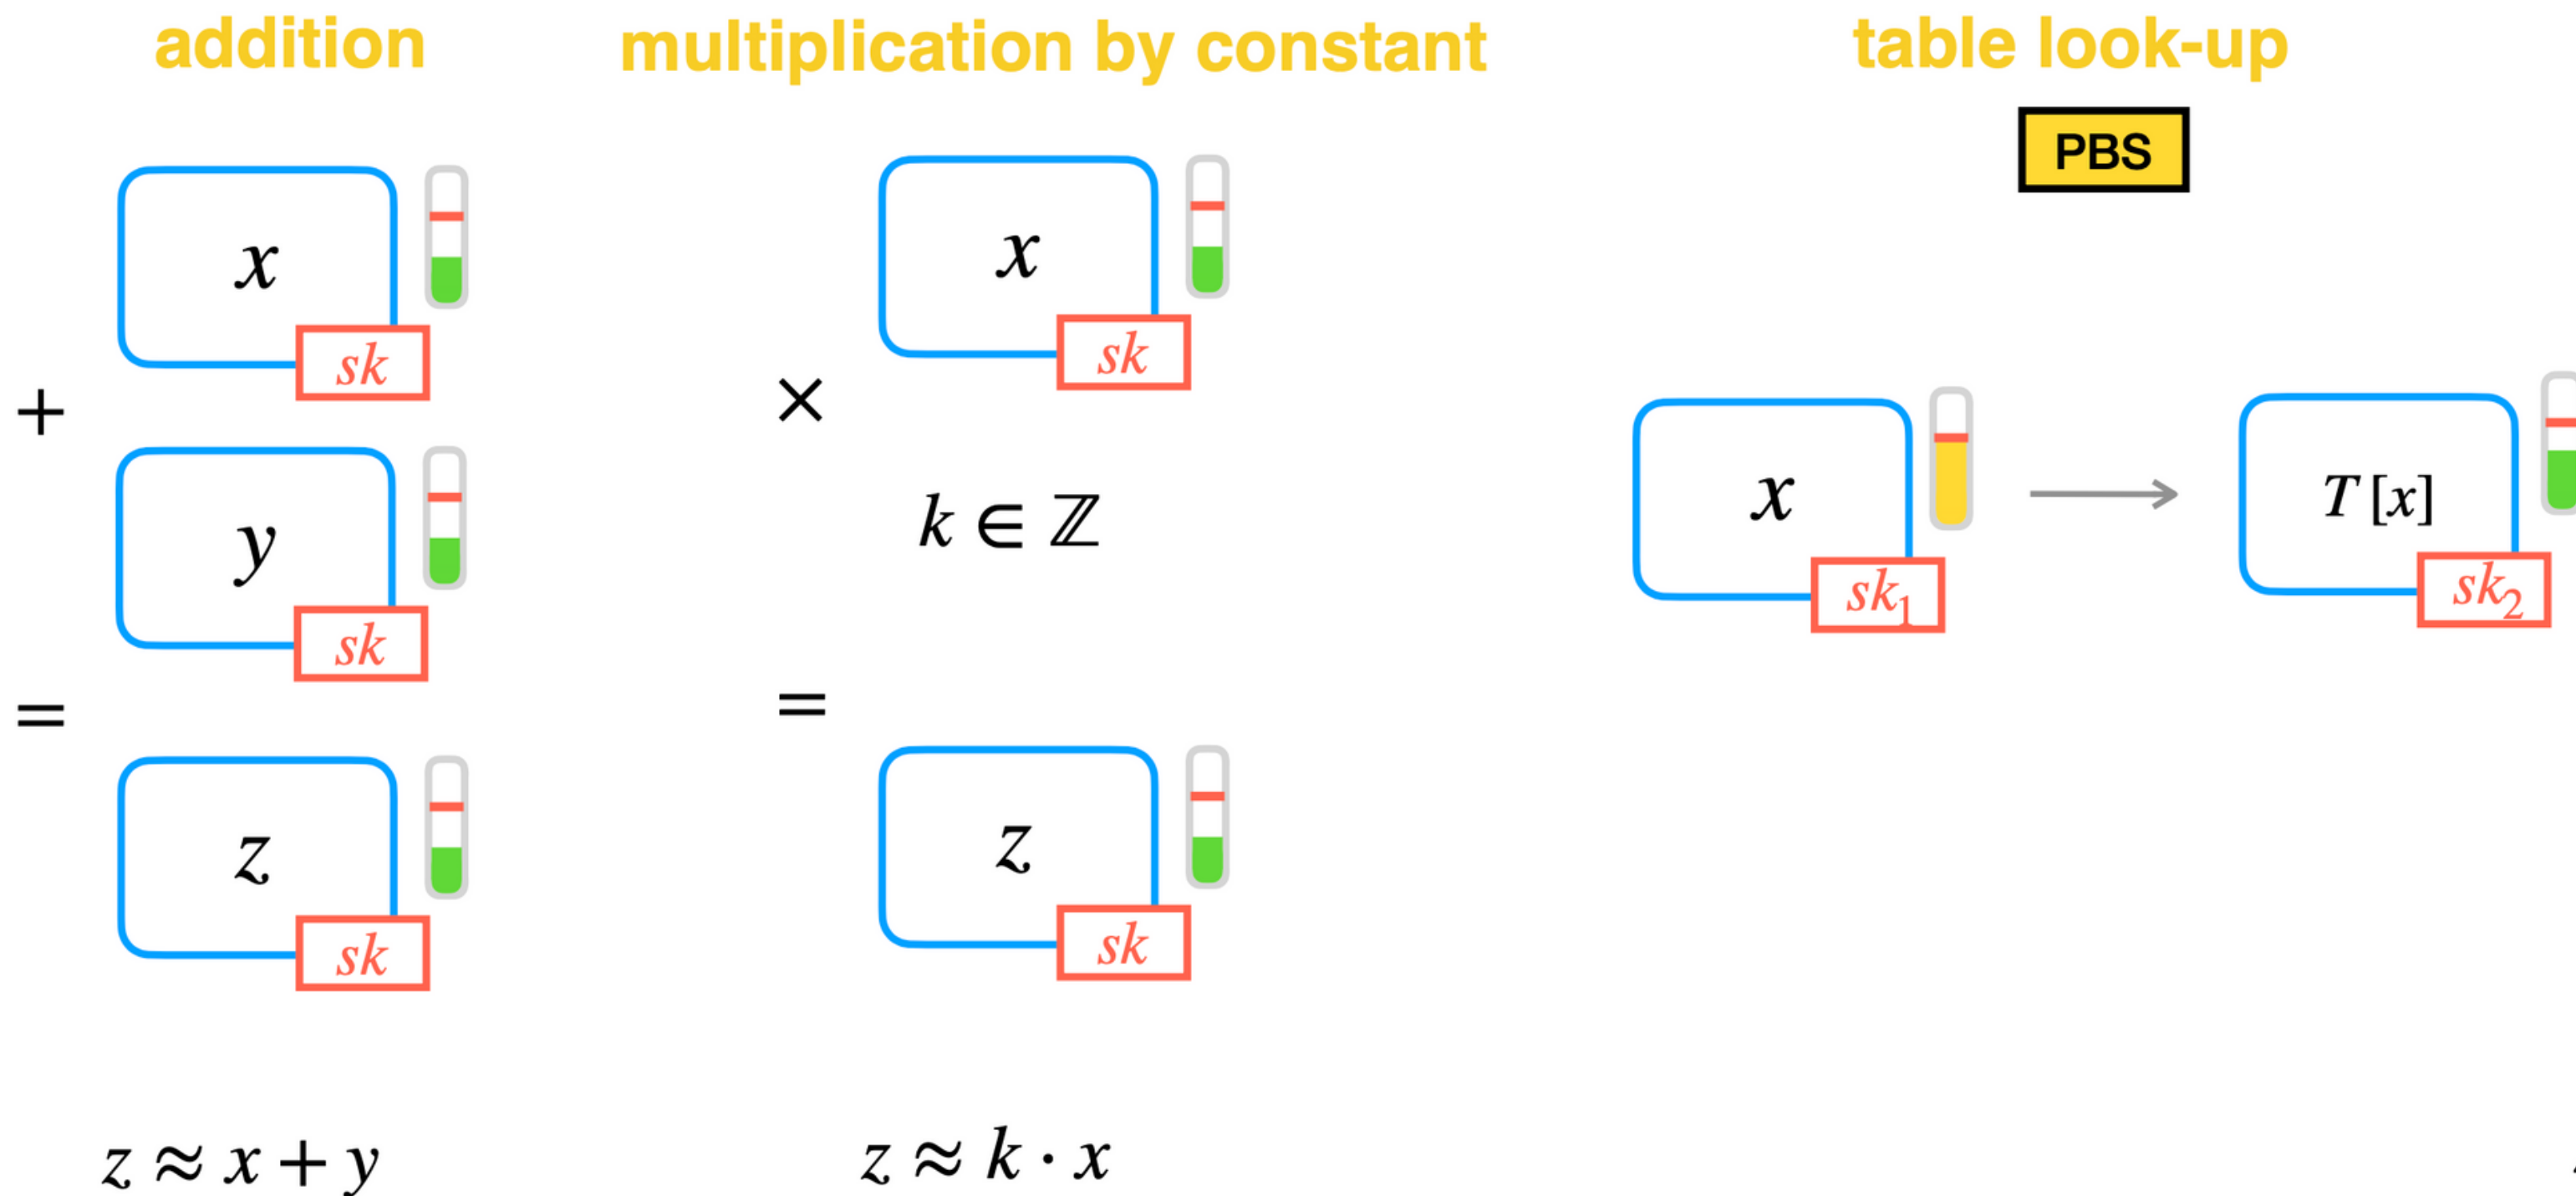


Privacy-Preserving Tree-Based Inference with Fully Homomorphic Encryption

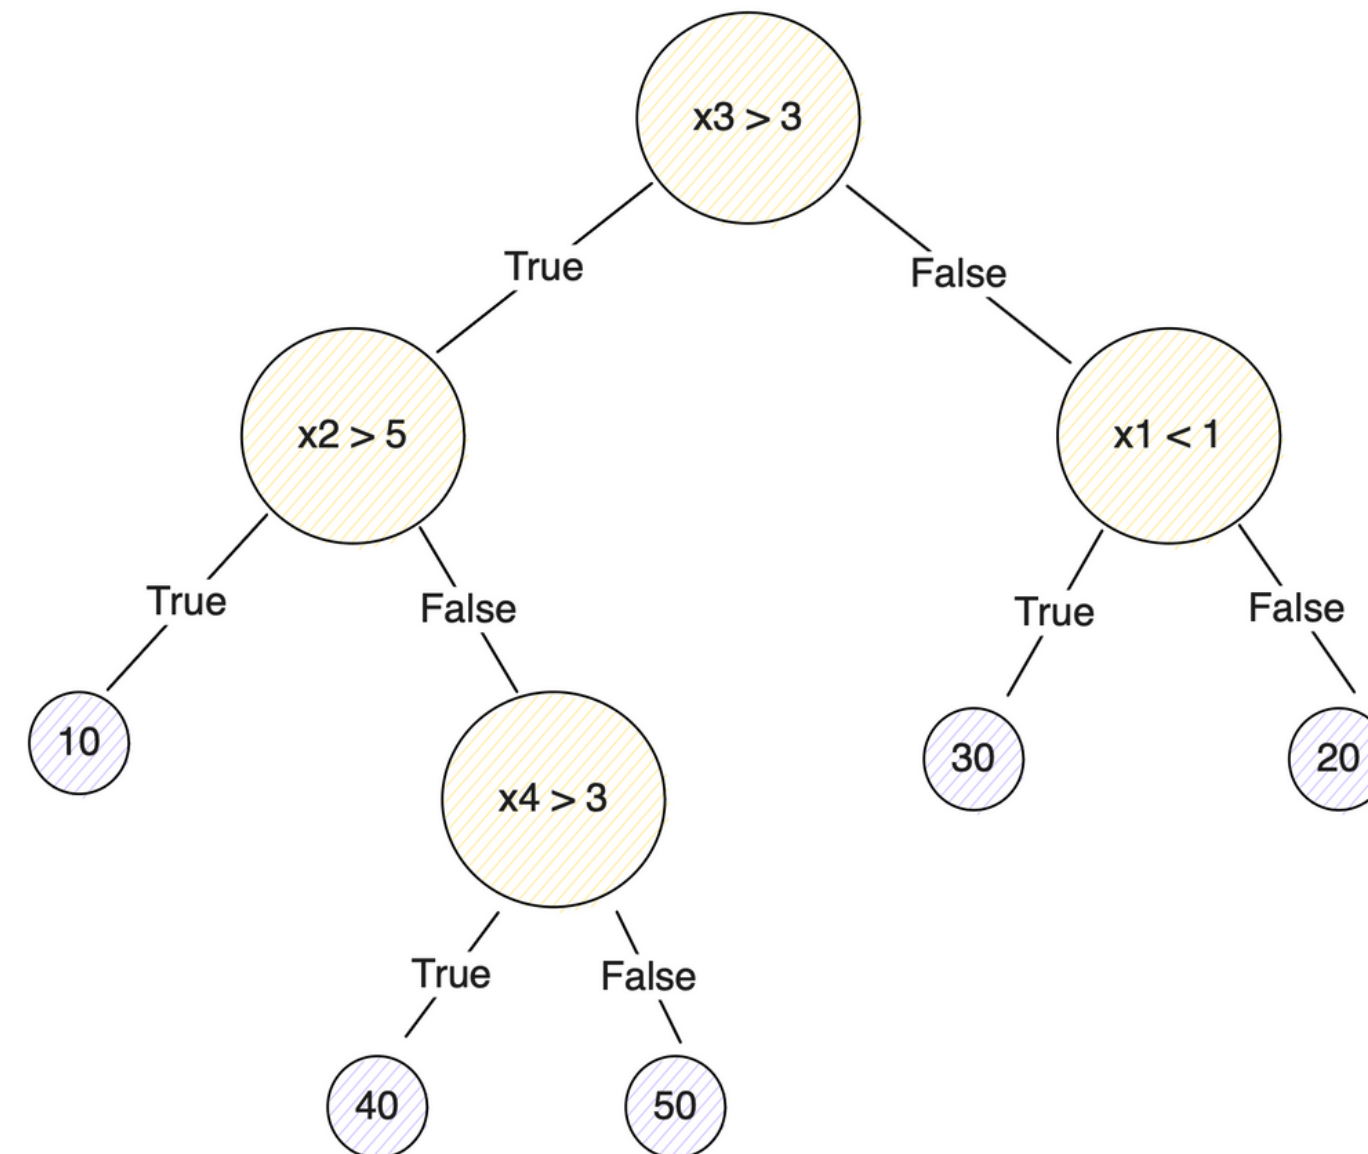
J. Frery, A. Stoian, R. Bredehoft, L. Montero, C. Kherfallah, B. Chevallier-Mames* and A. Meyre

Torus Fully Homomorphic Encryption



Doing Trees in FHE?

- Powerful and highly used ML models (eg, DecisionTree, RandomForest, XGBoost)
- Depends on conditions and branches, which are not doable in FHE



What We Describe in the Paper

A METHOD FOR TREES IN FHE

- Only matrix operations and two layer of PBS
- Adapt to any tree depths or number of trees
- Pretty efficient and versatile


AN IMPLEMENTATION

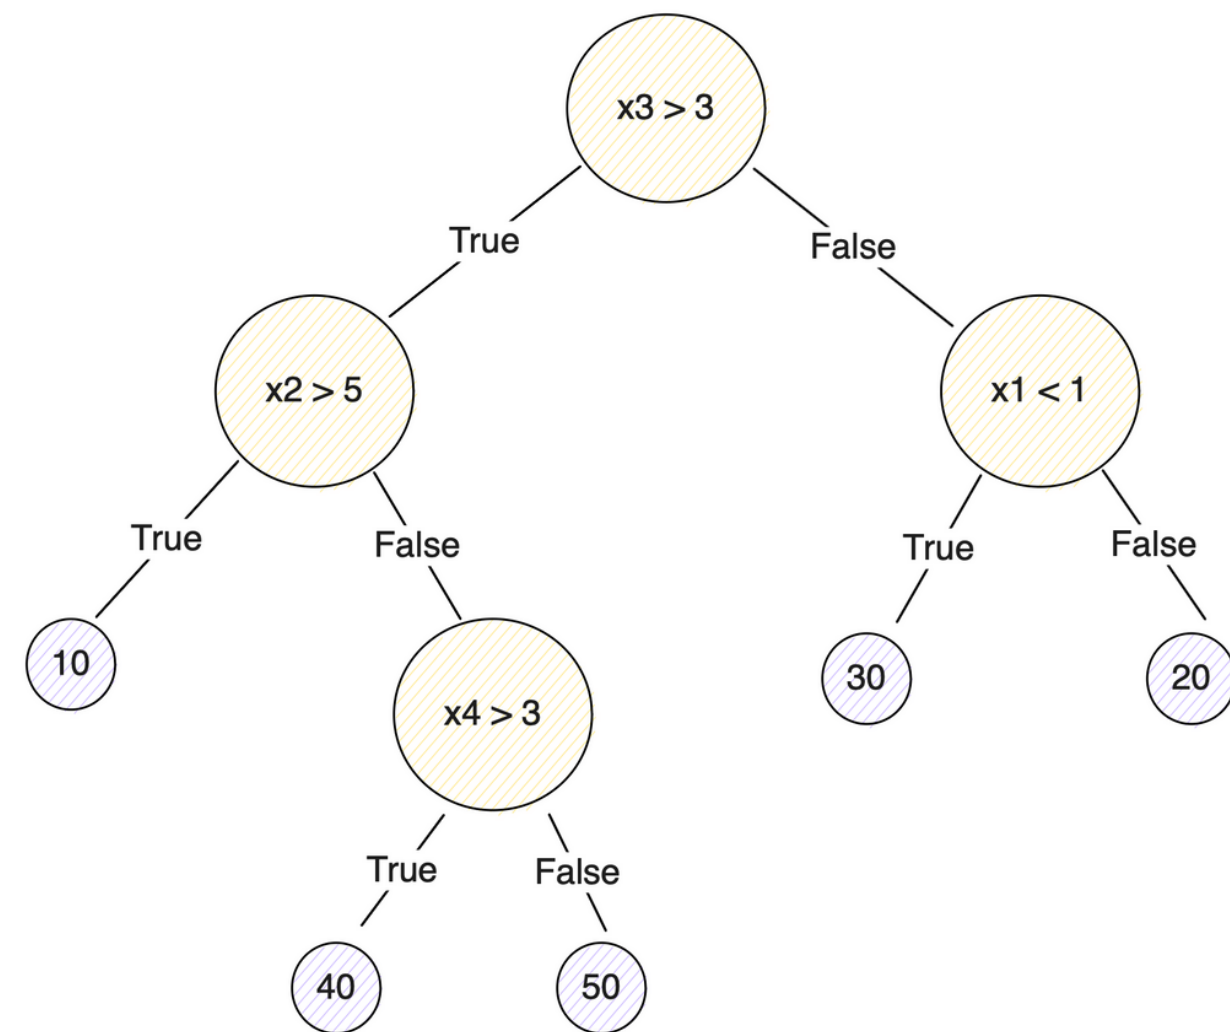
- Completely integrated in Concrete ML
- Open source
- Easy to use without crypto knowledge

EXPERIMENTAL RESULTS

- Experimental results over several datasets
- Available live demos on Hugging Face

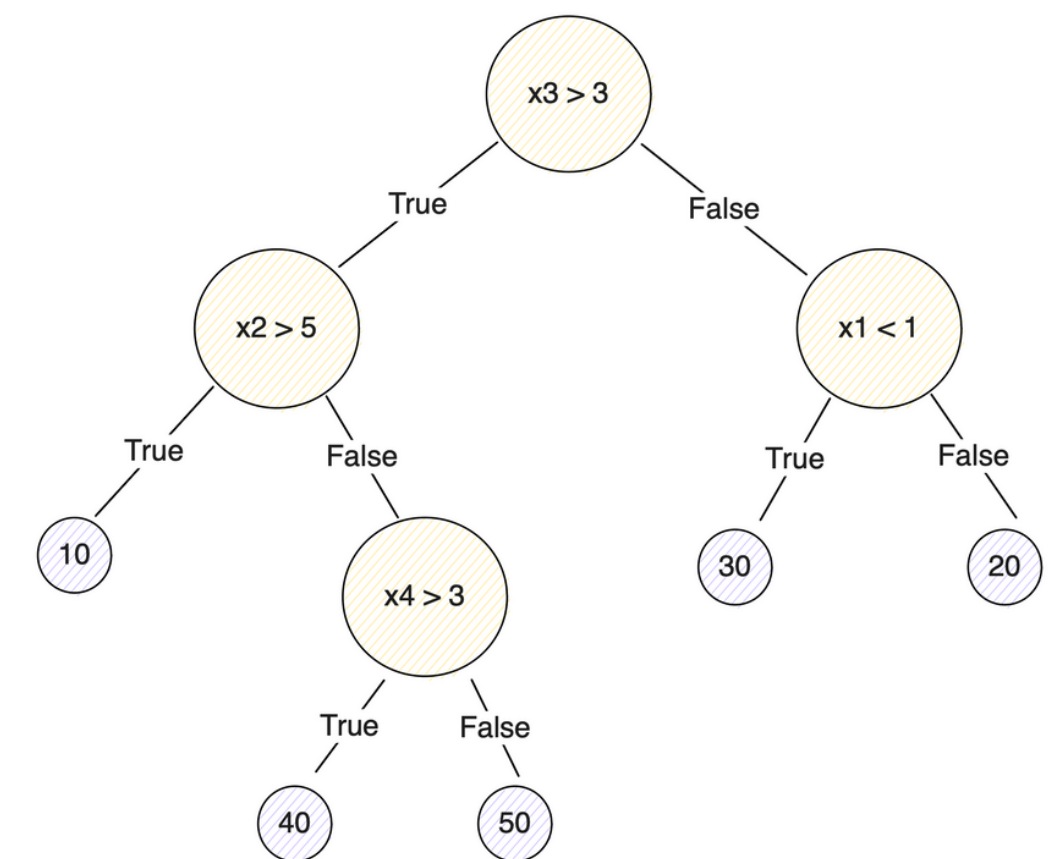
The Method

- Using **HummingBird** 
- Use TLU for $x_i > c$: eg, for $x > 2$, use $T[i] = [0, 0, 0, 1, 1, \dots, 1]$
- Have a **first layer of PBS** to compute the conditions
- Have a **second layer of PBS** to have a one-hot vector of which branch is taken. Eg, 10 will correspond to $[1, 0, 0, 0, 0]$ while 50 will correspond to $[0, 0, 1, 0, 0]$
- Have a final matrix multiplication to associate the one-hot with the final tree value

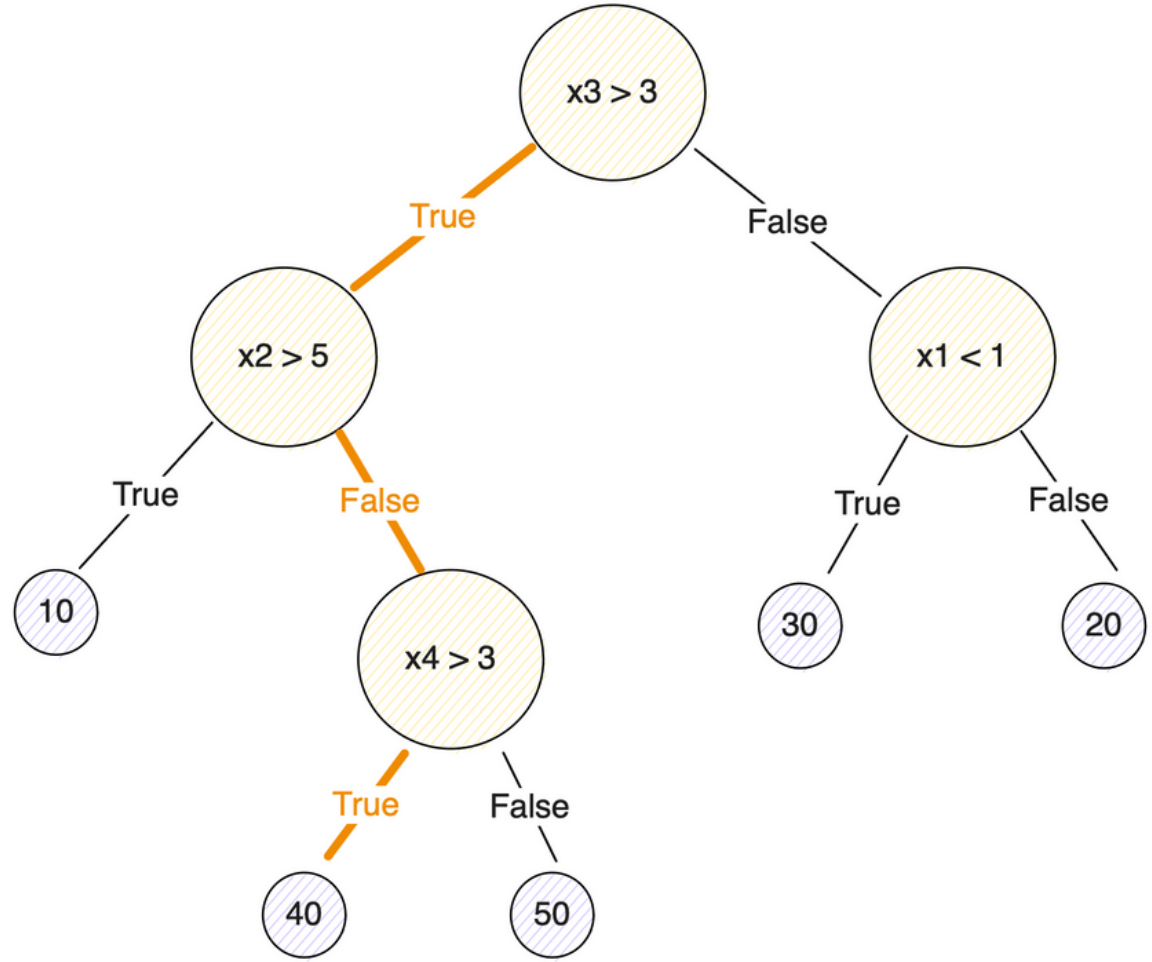



The Method – How To Have the One-Hot

- Eg, the 30 output, which has the one-hot encoding [0, 0, 0, 1, 0] appears iff:
 - $c1 := x3 > 3$ is False
 - $c3 := x1 < 1$ is True
 - whatever $c2$ and $c4$
- So, we say that the 4-th bit of the one-hot is: **$c3 - c1 == 1$**
- So, one linear layer to accumulate conditions, and one PBS to find if it's the "max" value



The Method



A

0	0	1	0
0	1	0	0
1	0	0	0
0	0	0	1
0	0	0	0

B

3	5	1	3
---	---	---	---

C


1	1	1	-1	-1
1	-1	-1	0	0
0	0	0	1	-1
0	1	-1	0	0

D

2	2	1	1	0
---	---	---	---	---

E

10
40
50
30
20



P = Input * A

Q = Conditions

R = Conditions * C

S = One-hot branch

T = Result

1	2	7	5	4
---	---	---	---	---

7	2	1	5
---	---	---	---

1	0	0	1
---	---	---	---

1	2	0	-1	-1
---	---	---	----	----

0	1	0	0	0
---	---	---	---	---

40

Implementation




```
from concrete.ml.sklearn import XGBClassifier


model = XGBClassifier(n_bits=8)
model.fit(X_train, y_train)
model.predict(X_test)
model.compile(X_train)
model.predict(X_test, fhe="simulate")
model.predict(X_test, fhe="execute")
```


Experimental Results

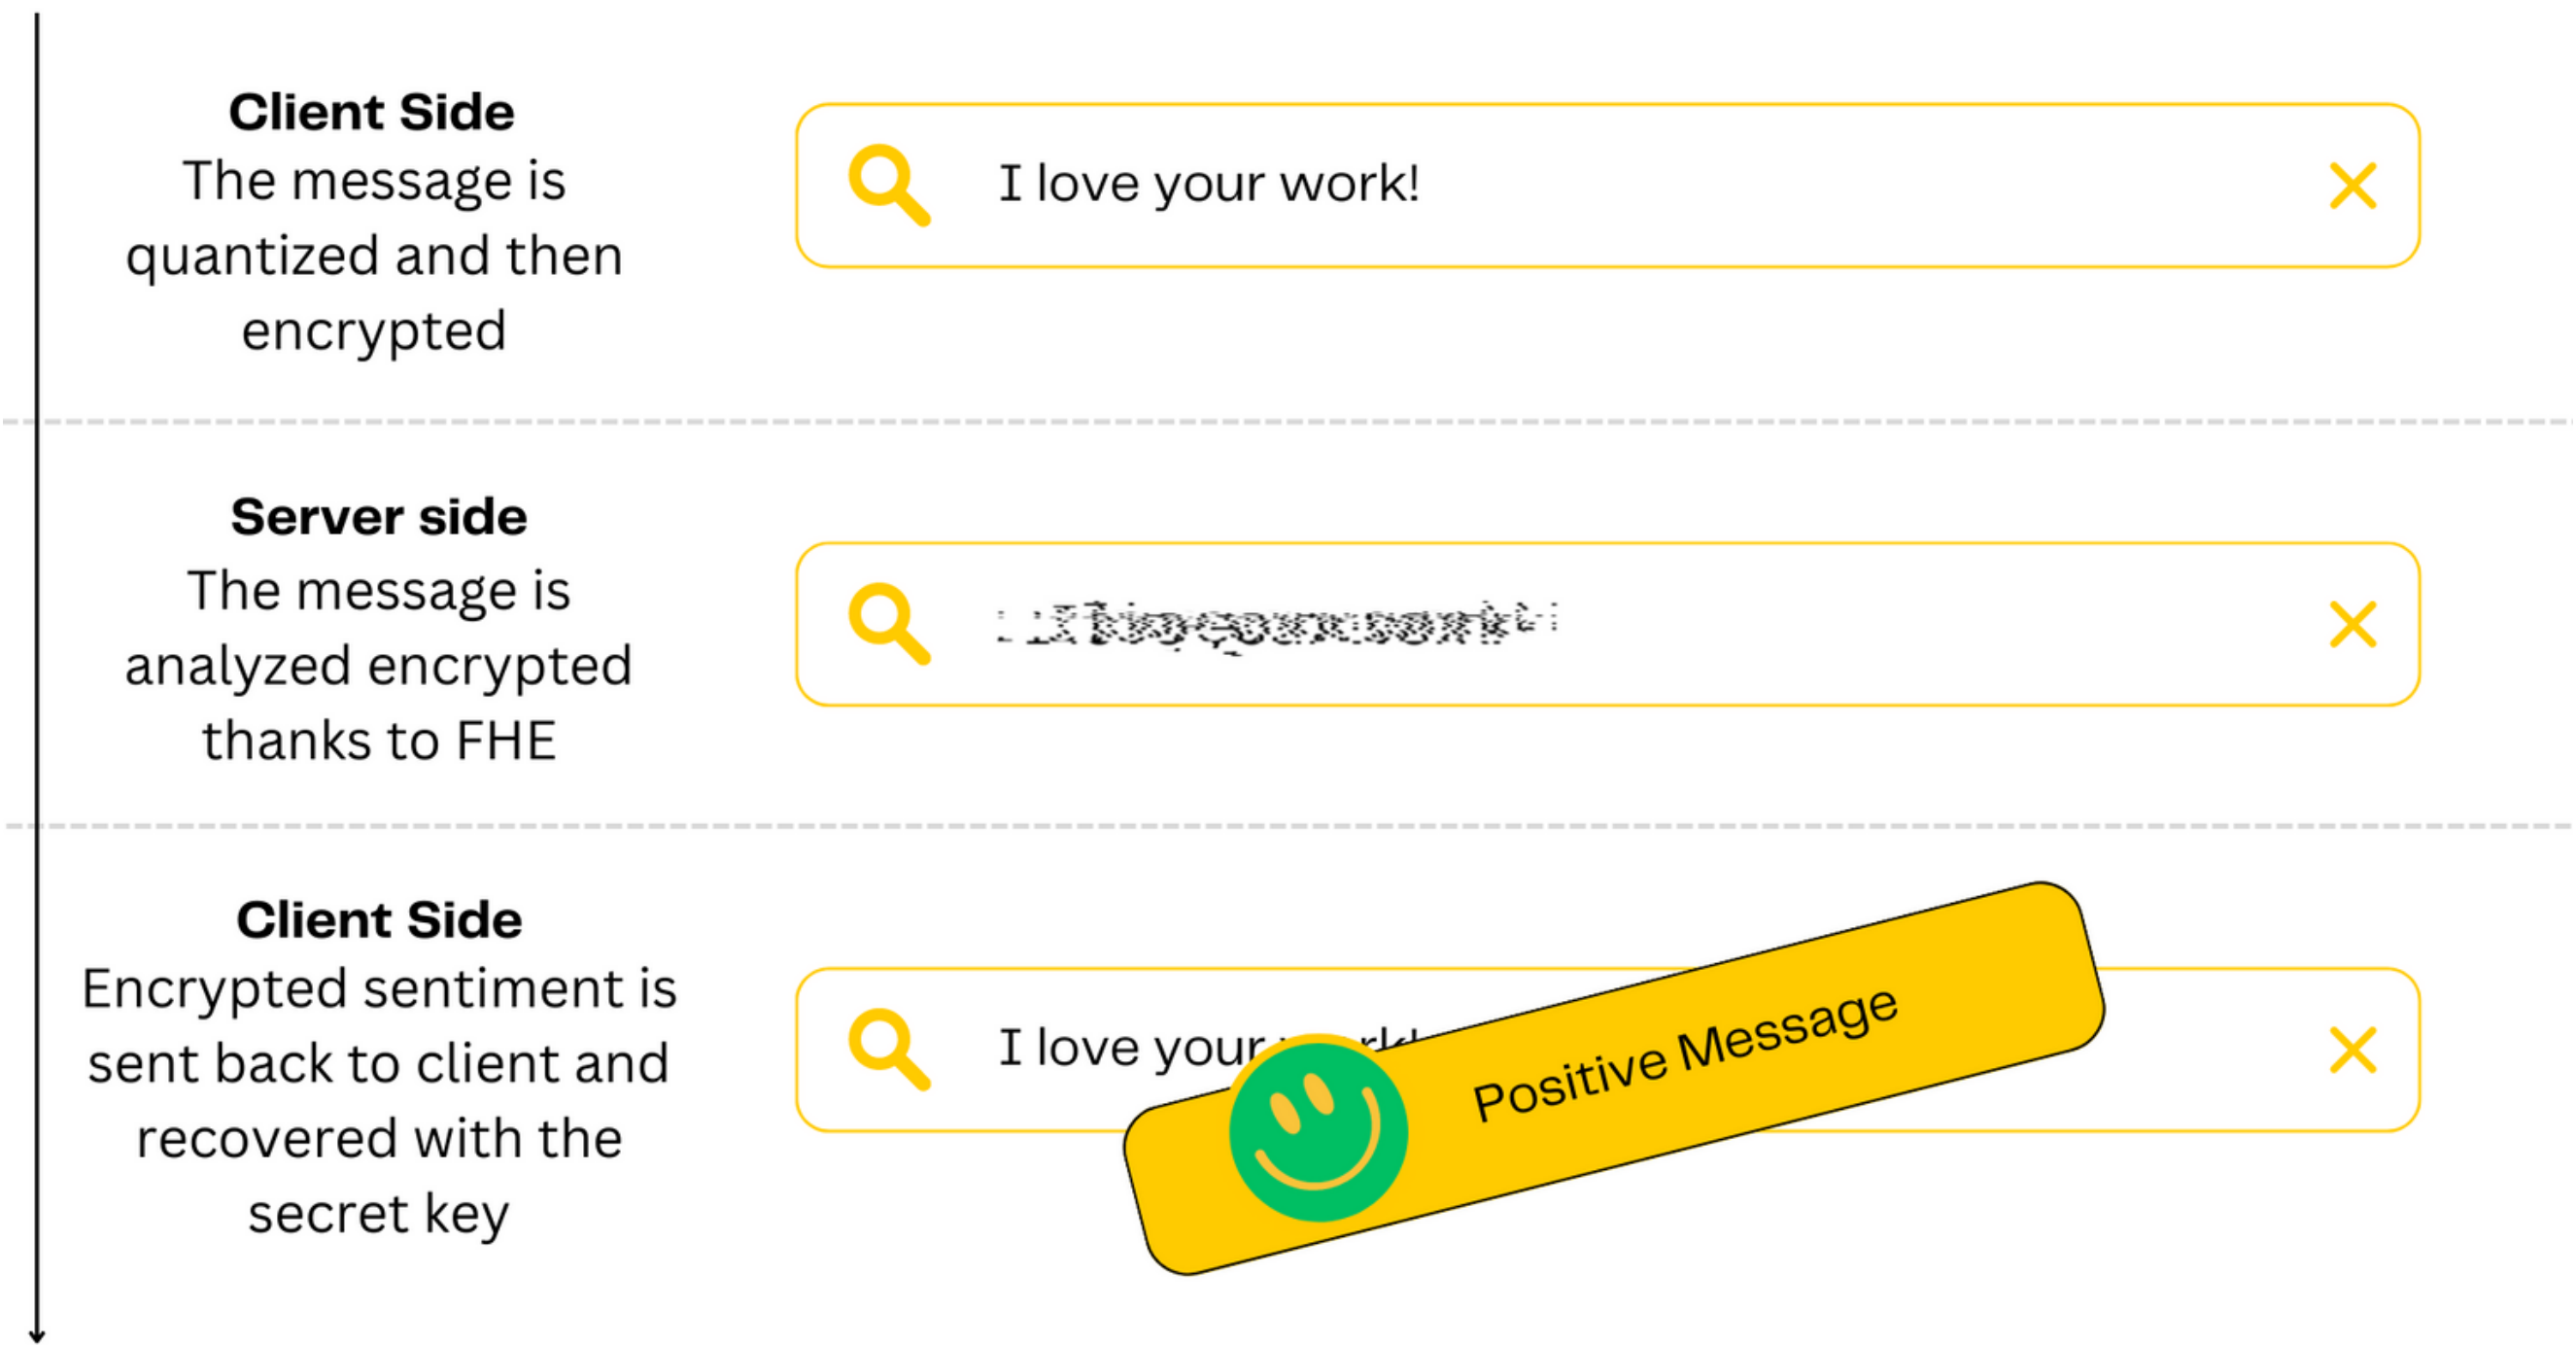
Inference time for different quantization bits for Concrete XGBoost Model



Accuracy for different quantization bits for Concrete XGBoost Model



Live Demo



Contact and links

- <https://eprint.iacr.org/2023/258.pdf>
- zama.ai
- github.com/zama-ai/concrete-ml
- <https://huggingface.co/zama-fhe>
- discord.fhe.org

Thank you !

